

# Pre-processing measures

## Imports

```
library(fairmodels)
library(tidymodels)
```

```
## -- Attaching packages ----- tidymodels 0.2.0 --
```

```
## v broom      0.8.0    v recipes      0.2.0
## v dials      0.1.1    v rsample      0.1.1
## v dplyr      1.0.9    v tibble      3.1.7
## v ggplot2    3.3.6    v tidyr       1.2.0
## v infer      1.0.0    v tune        0.2.0
## v modeldata  0.1.1    v workflows   0.2.6
## v parsnip    0.2.1    v workflowsets 0.2.1
## v purrr      0.3.4    v yardstick   0.0.9
```

```
## -- Conflicts ----- tidymodels_conflicts() --
```

```
## x purrr::discard() masks scales::discard()
## x dplyr::filter()  masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x recipes::step()  masks stats::step()
## * Learn how to get started at https://www.tidymodels.org/start/
```

```
library(rpart)
```

```
##
## Attaching package: 'rpart'

## The following object is masked from 'package:dials':
##
##      prune
```

```
library(discrim)
```

```
##
## Attaching package: 'discrim'

## The following object is masked from 'package:dials':
##
##      smoothness
```

```
source("../scripts/metrics_on_dataset.R")

## -- Attaching packages ----- tidyverse 1.3.1 --

## v readr 2.1.2      v forcats 0.5.1
## v stringr 1.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x readr::col_factor() masks scales::col_factor()
## x purrr::discard()     masks scales::discard()
## x dplyr::filter()      masks stats::filter()
## x stringr::fixed()     masks recipes::fixed()
## x dplyr::lag()          masks stats::lag()
## x readr::spec()        masks yardstick::spec()
```

## Data

```
df <- read_rds("../data/selection.rds") %>%
  select(-rating, -gender)
```

## Massaging

```
nB_df <- df %>%
  mutate(accepted = as.factor(accepted),
         index = 1:1000)
fitted_Bayes_model <- naive_Bayes() %>%
  fit(accepted ~ nationality + test_score + english_cert + extracurricular, nB_df)
nB_predictions <- predict(fitted_Bayes_model, nB_df, type="prob")
nB_df <- nB_df %>%
  mutate(prob_FALSE = nB_predictions[[".pred_FALSE"]],
         prob_TRUE = nB_predictions[[".pred_TRUE"]])
prob_true = sum(nB_df$accepted == "TRUE")/nrow(nB_df)
non_dutch_acceptance = nrow(filter(nB_df, accepted == "TRUE", nationality == "Non_Dutch")) /
  nrow(filter(nB_df, nationality == "Non_Dutch"))
print(non_dutch_acceptance)
```

```
## [1] 0.111332
```

```
while(non_dutch_acceptance < round(prob_true, 3)) {
  # Flip highest Non-Dutch
  highest_non_dutch <- nB_df %>%
    filter(accepted == "FALSE", nationality == "Non_Dutch") %>%
    arrange(desc(prob_TRUE)) %>%
    slice_head(n = 1) %>%
    select(index) %>%
    as.integer()
```

```

# Flip lowest Dutch
lowest_dutch <- nB_df %>%
  filter(accepted == "TRUE", nationality == "Dutch") %>%
  arrange(prob_TRUE) %>%
  slice_head(n = 1) %>%
  select(index) %>%
  as.integer()

nB_df <- nB_df %>%
  mutate(accepted = replace(accepted, index == highest_non_dutch, "TRUE"),
         accepted = replace(accepted, index == lowest_dutch, "FALSE"))

non_dutch_acceptance = nrow(filter(nB_df, accepted == "TRUE", nationality == "Non_Dutch")) /
  nrow(filter(nB_df, nationality == "Non_Dutch"))
# print(non_dutch_acceptance)
}
nB_df <- nB_df %>%
  select(-index, -prob_FALSE, -prob_TRUE)

model <- decision_tree(mode = "classification")
massaging_results <- all_metrics(nB_df, model, df)

```

```

## 'summarise()' has grouped output by 'nationality'. You can override using the
## '.groups' argument.
## 'summarise()' has grouped output by 'nationality'. You can override using the
## '.groups' argument.

```

```

print_all_metrics("Massaging", massaging_results)

```

```

## [1] "Massaging"
## [1] "Group fairness"
## # A tibble: 4 x 4
##   nationality accepted total  perc
##   <fct>         <fct>    <int> <dbl>
## 1 Dutch        FALSE     416  83.7
## 2 Dutch        TRUE       81  16.3
## 3 Non_Dutch    FALSE     421  83.7
## 4 Non_Dutch    TRUE       82  16.3
## [1] "Causal discrimination"
## [1] 0
## [1] "Unawareness"
## # A tibble: 4 x 4
##   nationality predicted_accepted total  perc
##   <fct>         <fct>          <int> <dbl>
## 1 Dutch        FALSE           416  83.7
## 2 Dutch        TRUE            81  16.3
## 3 Non_Dutch    FALSE           409  81.3
## 4 Non_Dutch    TRUE            94  18.7

```

## Reweighting

```
weights <- reweight(df$nationality, as.numeric(df$accepted))
weighted_model <- rpart(accepted ~ nationality + test_score + english_cert + extracurricular,
  df, weights)
weighted_df <- df %>%
  mutate(predicted = as.logical(predict(weighted_model, df)))
print("Reweightings")
```

```
## [1] "Reweightings"
```

```
print("Group fairness")
```

```
## [1] "Group fairness"
```

```
print(group_fairness(weighted_df, nationality, predicted)[[1]])
```

```
## 'summarise()' has grouped output by 'nationality'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 4 x 4
##   nationality predicted total  perc
##   <fct>         <lgl>    <int> <dbl>
## 1 Dutch        FALSE     390  78.5
## 2 Dutch        TRUE      107  21.5
## 3 Non_Dutch    FALSE     447  88.9
## 4 Non_Dutch    TRUE       56  11.1
```

```
print("Causal Discrimination")
```

```
## [1] "Causal Discrimination"
```

```
pop_size <- nrow(df)
# Flip nationalities
inverted_df <- df %>%
  mutate(nationality = ifelse(nationality == "Dutch", "Non_Dutch", "Dutch"))
# Add prediction column
eval_df <- df %>%
  mutate(inverted_accepted = as.logical(predict(weighted_model, inverted_df)),
    different = accepted != inverted_accepted)
print(sum(eval_df$different)/pop_size)
```

```
## [1] 0.123
```

```
print("Unawareness")
```

```
## [1] "Unawareness"
```

```
weighted_unaware_model <- rpart(accepted ~ test_score + english_cert + extracurricular, df, weights)
weighted_unaware_df <- df %>%
  mutate(predicted = as.logical(predict(weighted_unaware_model, df)))
print(group_fairness(weighted_unaware_df, nationality, predicted)[[1]])
```

```
## 'summarise()' has grouped output by 'nationality'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 4 x 4
##   nationality predicted total perc
##   <fct>         <lgl>    <int> <dbl>
## 1 Dutch        FALSE     390  78.5
## 2 Dutch        TRUE      107  21.5
## 3 Non_Dutch    FALSE     379  75.3
## 4 Non_Dutch    TRUE      124  24.7
```

## Resampling

```
resample_indexes <- resample(df$nationality, as.numeric(df$accepted))
resampled_df <- df[resample_indexes, ]
model <- decision_tree(mode = "classification")
resampling_results <- all_metrics(resampled_df, model, df)
```

```
## 'summarise()' has grouped output by 'nationality'. You can override using the
## '.groups' argument.
## 'summarise()' has grouped output by 'nationality'. You can override using the
## '.groups' argument.
```

```
print_all_metrics("Resampling", resampling_results)
```

```
## [1] "Resampling"
## [1] "Group fairness"
## # A tibble: 4 x 4
##   nationality accepted total perc
##   <fct>         <lgl>    <int> <dbl>
## 1 Dutch        FALSE     416  83.7
## 2 Dutch        TRUE       81  16.3
## 3 Non_Dutch    FALSE     421  83.7
## 4 Non_Dutch    TRUE       82  16.3
## [1] "Causal discrimination"
## [1] 0.123
## [1] "Unawareness"
## # A tibble: 4 x 4
##   nationality predicted_accepted total perc
##   <fct>         <fct>          <int> <dbl>
## 1 Dutch        FALSE             402  80.9
## 2 Dutch        TRUE              95  19.1
## 3 Non_Dutch    FALSE            435  86.5
## 4 Non_Dutch    TRUE             68  13.5
```